

Information Retrieval with PostgreSQL

Alexander Hebel

Heidelberg University
Institute of Computer Science
Database Systems Research Group
vx228@uni-heidelberg.de

Mai 6, 2020

Outline

- 1 Introduction
- 2 Approach and Realizations
- 3 Adding Functionality to Tsvectors
- 4 Evaluation Objectives
- 5 Conclusion

Outline

- 1 Introduction
- 2 Approach and Realizations
- 3 Adding Functionality to Tsvectors
- 4 Evaluation Objectives
- 5 Conclusion

Information Retrieval System (IRS)

Information Retrieval (IR)

- Information retrieval is the science of searching for information in a document
- An IR system is a software system that provides access to books, journals and other documents; stores and manages those documents.
- Web search engines are the most visible IR applications.

Full-Text-Search

The activity of searching through a collection of natural-language documents to locate those that best match a query

Objectives

To our best knowledge there exists no IRS based on a relational database.

- Finding different suitable database models
- Python api for the database creation and communication
- Crawl Wiki pages to gather text data
- Store and manage those text documents
- Support search queries containing the boolean operator AND
- Create a ranking function for the query results
- Compare the ranking of whole Wiki pages and their sections

Tsvector and Tsquery

- PostgreSQL provides **two** data types that are designed to support full text search
- The **tsvector** type represents a document in a form optimized for text search
- The **tsquery** type similarly represents a text query

Example Tsvector

```
SELECT to_tsvector('english', 'The Fat Rats');  
{'fat':2 'rat':3}
```

Example Tsquery

```
SELECT to_tsquery('Fat & Cats');  
{'fat' & 'cat'}
```

Tsvector Pro and Contra

Possibilities

- Full text search
- GIN-Index
- Automatic tokenization and lemmatization
- Adding weights
- Predefined ranking function

Limitations

- The number of lexemes must be less than 2^{64}
- Max position value: 16383
- No more than 256 positions per lexeme
- Relative small set of manipulation methods
- Ranking function must be customized

Outline

- 1 Introduction
- 2 Approach and Realizations**
- 3 Adding Functionality to Tsvectors
- 4 Evaluation Objectives
- 5 Conclusion

Text Data Statistics


Document Definition

A document in my case is either a page title, section title or section text

Wikipedia Category "Sports"

- Number of Wiki pages: 2000
 - Number of sections (captions also count as section): 45756
 - Total number of lexemes: 1969427
-
- Max number of words of a page: 124136
 - Max number of lexemes of a page: **17787**
-
- Max number of words of a section: 24163
 - Max number of lexemes of a section: **2785**

Database Model



Outline

- 1 Introduction
- 2 Approach and Realizations
- 3 Adding Functionality to Tsvectors**
- 4 Evaluation Objectives
- 5 Conclusion

Reason to Extend PostgreSQL

Only a few manipulation methods for tsvector exist

For example it is not possible to:

- count the occurrence of **all** lexemes in a tsvector (only the number of distinct lexemes)
- create a tsvector with an offset for the position index
- get the max index of a tsvector
- concatenate two tsvector without changing the position values

Solution

Write your own function either in plpgsql language or in C

Custom C-Functions in PostgreSQL

Prerequisites

- Developer version of PostgreSQL
- Installation of make
- Root privilege on database

Folder structure

Extension

- ─ function.c
- ─ Makefile
- ─ function.control
- ─ function--1.0.sql
- ─ README.function

Steps

- (1) make install
- (2) CREATE EXTENSION "extension"

Outline

- 1 Introduction
- 2 Approach and Realizations
- 3 Adding Functionality to Tsvectors
- 4 Evaluation Objectives**
- 5 Conclusion

Evaluation Objectives

- Three different database models have been tested
- The model mainly using tsvector has proven most suitable
 - + GIN index for performance
 - + ts_rank() function for ranking
 - + ts_query boolean and even phraseto_tsquery() support


Calculation of Ranking

- Function ts_rank() of PostgreSQL
- Considers word occurrences and distance between words

Relationship Between Page and Section Ranking

- **section ranking:** $\text{ranking} / \text{num_words_of_section}$
- **page ranking:** $\text{sum_of_rankings} / \text{num_words_of_page}$

Query:"game", Sorted by Max Section Ranking



Distance Between Rankings

Calculating Distance

- Two sorted lists (section_rankings and page_rankings)
- Distance from first section in section_rankings to position of page in page_rankings containing this section and so on...

Query "game"

- result pages: 1176
- Top 10 - 289.5 avg(dist)
- Top 20 - 250.9 avg(dist)
- Top 30 - 183.2 avg(dist)
- Top 40 - 155.5 avg(dist)

Query "game & team & ball"

- result pages: 274
- Top 10 - 36.0 avg(dist)
- Top 20 - 40.6 avg(dist)
- Top 30 - 36.6 avg(dist)
- Top 40 - 34.0 avg(dist)

Outline

- 1 Introduction
- 2 Approach and Realizations
- 3 Adding Functionality to Tsvectors
- 4 Evaluation Objectives
- 5 Conclusion**

Conclusion and Future Work

Conclusion

- An IR system made of a relational database is possible
- Rankings for sections and page return total different results
- Tsvector has a lot of potential
- PostgreSQL is easy customizable
- Solr gives a little bit different ranking results

Future Work

- Improve the ranking algorithm with tf idf information (ts_stat)
- Boost certain terms of a query
- Proximity query (aka sloppy phrase query)
- Tests on big datasets

Questions

Questions

Query:"game", Sorted by Sum of Section Rankings

